



University of Stuttgart
Institute of Parallel and Distributed Systems



**Distributed
Systems**

Fachpraktikum / Lab-Course

Software-Defined and Time-Sensitive Networking

Tutorial: Time Sensitive Networking
Part 1: Background and TSN in Linux

Frank Dürr

**Summer
Term
2023**

Agenda

- Motivation
- TSN Background
- TSN in Linux

Motivation

Many cyber physical systems are safety critical!

- Failure to react to physical world in time leads to substantial damage or even harm of people
- ➔ Deadlines must be met deterministically

Deterministic Real-Time Communication required for networked real-time systems

➔ Deterministically bounded network delay


Motivation

Example

Printing and Folding Machine



Motors:
> 50 axes




Hochwertig bedruckte Verp
aus Wellpappe stellen imm
Anforderungen an die Druck

Final result: printed and folded box

[Source: Göpfer Maschinen GmbH]

Motivation


Old Concept: Mechanical Synchronization



[Source: Prof. Dr. -Ing. J. M. Pacas, Universität Siegen
Prof. Dr. -Ing. R. M. Kennel, Technische Universität München]

Motivation

New Concept: Electronic Motion Control



[Source: Prof. Dr. -Ing. J. M. Pacas, Universität Siegen
Prof. Dr. -Ing. R. M. Kennel, Technische Universität München]

Sensor and actuators:
servo motors connected to motion controller via **real-time communication network**

Motivation

Requirement: Low Deterministic Network Delay and Jitter

- **Cycle time:** time to read/update all sensors/actuators once
 - Defines update rate: each device once per cycle
 - Defines max. latency: all devices updated within same cycle (deadline = period)
- **At high machine speeds, low cycle time is crucial**
 - Printing machine: > 300 m/min material throughput speed
 - 1 ms cycle time means one update every 5 mm
- **Industrial real-time networks: cycle times down to 31.25 μ s**
 - Example: SERCOS III (Ethernet-based)
 - Cycle times from 31.25 μ s to 65 ms
 - 31.25 μ s with 8 devices (axes), 500 μ s – 1 ms with 100 devices (axes)
 - Jitter < 100 ns

Motivation

Converged Networks

- Initially, dedicated networks for real-time and non-real-time traffic
 - Separation of OT networks and IT networks
- Dedicated networks are costly
- Dedicated networks used precious space, weight, energy

➔ Converge OT and IT into one single network

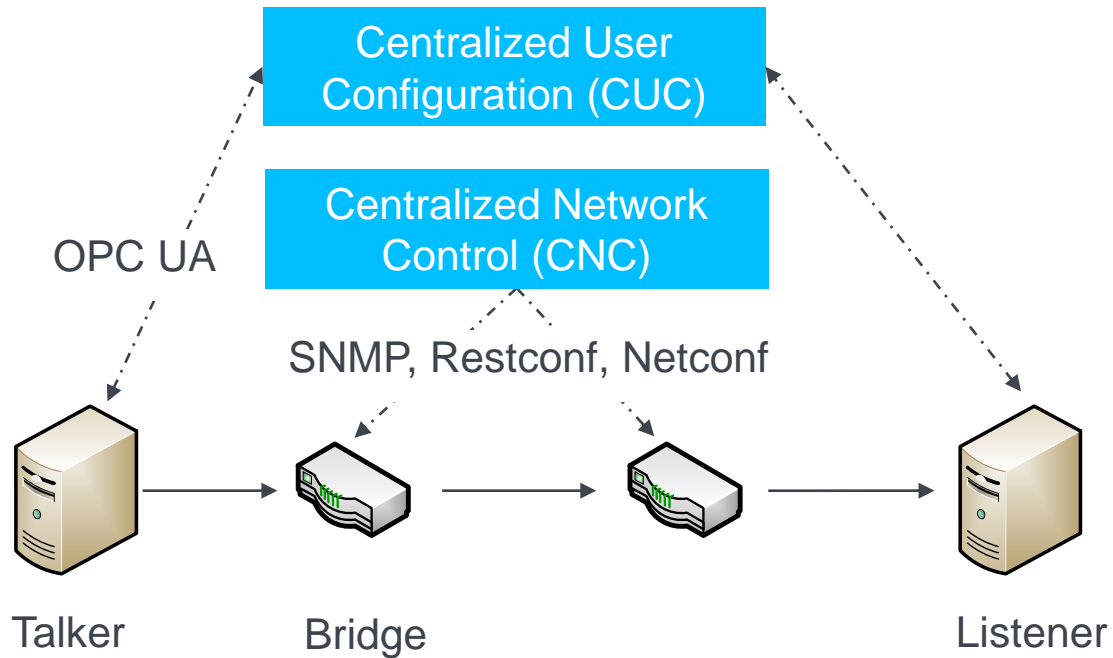
Agenda

- Motivation
- **TSN Background**
- TSN in Linux

Time Sensitive Networking

- Time Sensitive Networking (TSN) is one technology for converged networks based on Ethernet Standard
- IEEE Standard 802.1Q
- Enhanced popular Ethernet technology (IEEE 802.3)
 - Precise time synchronization: Precision Time Protocol (PTP)
 - Scheduling with very low, deterministic network delay and jitter: Time-Aware Shaper
 - Frame Replication and Elimination to increase reliability
 - Frame Preemption to let time-critical traffic preempt transmission of non-time critical traffic

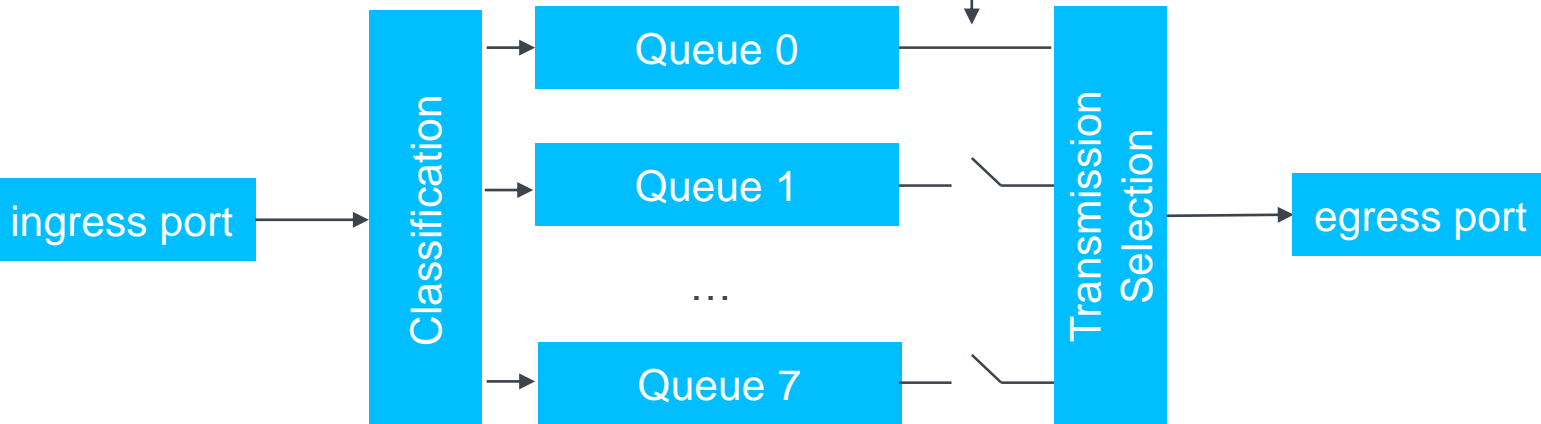
Architecture



Time Aware Shaper (IEEE 802.1 Qbv)

Gate Control List

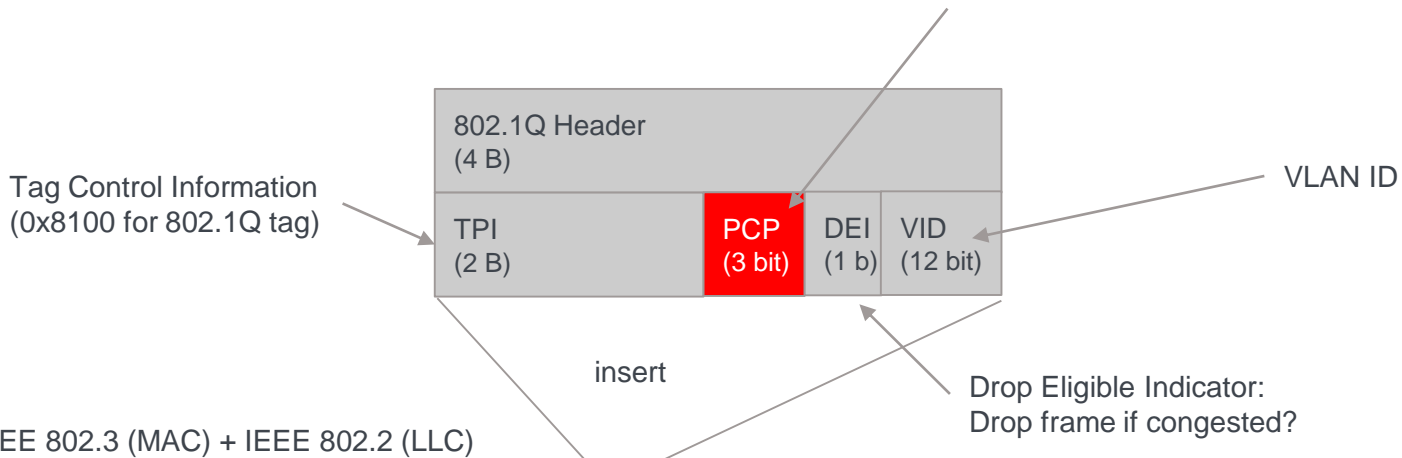
t	gate state
0 – 1000 ns	c o o o o o o o
1000 ns – 2000 ns	o c o o o o o o



Ethernet Frames

Definition of Priorities

Priority Code Point (PCP):
3 bits = 8 classes (see egress queues on previous slide)



Traffic Types

- Isochronous
 - Cyclic traffic with given period
 - Clocks of talkers and listeners are synchronized to working clock (network time as used by TSN bridges)
 - Transmission time of talkers can be aligned to gate control times
- Cyclic asynchronous
 - Cyclic traffic with given period
 - Clocks of talkers and listeners are not synchronized to working clock
 - Transmission cycles of talkers are not aligned to gate control times
- Sporadic / asynchronous
 - Aperiodic transmission

Agenda

- Motivation
- TSN Background
- **TSN in Linux**

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

Linux implements Time-Aware Shaping through a QDisc (Queueing Discipline): **TAPRIO (Time Aware Priority Shaper)**

- Controls egress traffic of network device
- Time-driven gates similar to Time-Aware Shaper
- Controlled through Command Line Interface
 - No control plane implementation (NETCONF, RESTCONF, SNMP)

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

Network
device

```
$ tc qdisc replace dev enp2s0f1 parent root handle 100
taprio \
num_tc 2 \
map 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 \
queues 1@0 1@1 \
base-time 1554445635681310809 \
sched-entry S 01 800000 sched-entry S 02 200000 \
clockid CLOCK_TAI
```

Number of
traffic classes

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

```
$ tc qdisc replace dev enp2s0f1 parent root handle 100
taprio \
num_tc 2 \
map 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 \
queues 1@0 1@1 \
base-time 1554445635681310809 \
sched-entry S 01 800000 sched-entry S 02 200000 \
clockid CLOCK_TAI
```

Mapping of
SKB priorities
0...15 to traffic
classes

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

```
$ tc qdisc replace dev enp2s0f1 parent root handle 100
taprio \
num_tc 2 \
map 1 0 1 1 1 1 1 1 1 1 1
queues 1@0 1@1 \
base-time 1554445635681310809 \
sched-entry S 01 800000 sched-entry S 02 200000 \
clockid CLOCK_TAI
```

Mapping of traffic classes to egress queues of device:

- 1st class: 1 queue at queue index 0
- 2nd class: 1 queue at queue index 1

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

```
$ tc qdisc replace dev enp2s0f1 parent root handle 100
taprio \
num_tc 2 \
map 1 0 1 1 1 1 1 1 1
queues 1@0 1@1 \
base-time 1554445635681310809 \
sched-entry S 01 800000 sched-entry S 02 200000 \
clockid CLOCK_TAI
```

Cycle base time in nano-seconds

System clock to use:
TAI (Temps Atomique International):
no leap seconds

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

```
$ tc qdisc replace dev taprio \
num_tc 2 \
map 1 0 1 1 1 1 1 1 1 \
queues 1@0 1@1 \
base-time 155444563568130809 \
sched-entry S 01 800000 sched-entry S 02 200000 \
clockid CLOCK_TAI
```

Gate Control List (Schedule):

- S 01: set gate state (only possibility)
- 01: gate 0 open (first bit)
- 02: gate 1 open (second bit)
- 800000: duration of entry in nano-secs

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

TAPRIO uses socket buffer (SKB) priorities for scheduling!

- Not directly the PCP
- Application can define priority for sockets:

```
setsockopt(  
    socket,  
    SOL_SOCKET,  
    SO_PRIORITY,  
    &priority, // int value  
    sizeof(priority));
```

- For egress and ingress packets:
PCP needs to be mapped to SKB priority and vice versa (see below)

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

TAPRIO requires network device to have 8 TX (transmit) queues

- Check number of TX queues of device:

```
$ ls /sys/class/net/enp2s0f1/queues/
```

```
rx-0 rx-1 rx-2 rx-3 rx-4 rx-5 rx-6 rx-7 tx-0 tx-1 tx-2  
tx-3 tx-4 tx-5 tx-6 tx-7
```

- Create virtual Ethernet device (veth) with several TX queues:

```
$ sudo ip link add vethname1 numtxqueues 8 type veth  
peer name vethname2 numtxqueues 8
```

Time-Aware Shaping in Linux

TAPRIO (Time Aware Priority Shaper)

Recommended reading:

- Man page of TAPRIO: `man taprio`

- Blog post:

https://www.frank-durr.de/posts/2019/04/11/software_tsn_switch.html

VLANs in Linux

Creating VLAN Devices

Sending a Ethernet frame with PCP requires a VLAN tag

- Creating a VLAN interface for an existing device (eth0):

```
$ sudo ip link add link eth0 \
```

```
name eth0.100 \
```


Name of VLAN device

```
type vlan id 100
```

VLAN ID

VLANs in Linux

Tagged and Untagged Frames



VLANs in Linux

PCP Mapping

SKB priorities can be mapped to PCP values for ingress and egress traffic:


```
$ sudo ip link set eth0.100 \  
type vlan \  
egress 0:0 1:1
```



Map SKB priority 1 to PCP 1

VLANs in Linux

Priority Mapping



Time and Clocks in POSIX

Getting Current Time

Implementing isochronous and cyclic talkers requires clocks:

Getting current time:

```
clock_gettime(CLOCK_MONOTONIC, &t);
```

- `CLOCK_MONOTONIC`: “A nonsettable system-wide clock that represents monotonic time since—as described by POSIX—“some unspecified point in the past”. [man page of `clock_gettime`]
- `CLOCK_TAI`: “A nonsettable system-wide clock derived from wall-clock time but ignoring leap seconds. This clock does not experience discontinuities and backwards jumps caused by NTP inserting leap seconds as `CLOCK_REALTIME` does.” [man page of `clock_gettime`]

Time and Clocks in POSIX

Sleeping

Sleeping for an interval or until an absolute point in time:

```
int clock_nanosleep(  
    clockid_t clockid, // Clocks see pervious slide  
    int flags,          // TIMER_ABSTIME to sleep until  
                        // a point in time  
    const struct timespec *request, // relative or  
                                    // absolute time  
    struct timespec *remain); // remaining time if  
                              // sleep is interrupted
```

Questions?